清华大学

Tsinghua University

算法分析与设计课程作业

主讲教师：王振波

姓名： <u>冯寅潇</u>

学号： <u>2020311833</u>

学院： <u>交叉信息研究院</u>

Algorithms Analysis and Design

Homework 1

p22, Ex1.1

Decide whether you think the following statement is true or false. If it is true, give a short explanation. If it is false, give a counterexample.

True or false? In every instance of the Stable Matching Problem, there is a stable matching containing a pair (m, w) such that m is ranked first on the preference list of w and w is ranked first on the preference list of m.

Answer

False. Counterexample is show blew:

Name	Most like	Least like
A	**X**	Y
B	**Y**	X
X	B	**A**
Y	A	**B**

The stable matching is $\{A - X, B - Y\}$. We can see that A and B are not ranked first on the preference list of X and Y. So, it is not necessary to have such pair in stable matching. Actually, there is no such pair at all in this example.

p22, Ex1.2

Decide whether you think the following statement is true or false. If it is true, give a short explanation. If it is false, give a counterexample.

True or false? Consider an instance of the Stable Matching Problem in which there exists a man m and a woman w such that m is ranked first on the preference list of w and w is ranked first on the preference list of m. Then in every stable matching S for this instance, the pair (m, w) belongs to S.

Answer

True. By contradiction, suppose there exists a stable matching S, the adequate pair $(m, w) \notin S$. Assume that a and b are assigned partners of m and w in this matching S. Then, m prefers w to its assigned partner and w prefers m to its assigned partner. So this pair is an unstable pair which contradicts the hypothesis. So, the statement is true.

p22, Ex1.4

Gale and Shapley published their paper on the Stable Matching Problem in 1962; but a version of their algorithm had already been in use for ten years by the National Resident Matching Program, for the problem of assigning medical residents to hospitals.

Basically, the situation was the following. There were m hospitals, each with a certain number of available positions for hiring residents. There were n medical students graduating in a given year, each interested in joining one of the hospitals. Each hospital had a ranking of the students in order of preference, and each student had a ranking of the hospitals in order of preference. We will assume that there were more students graduating than there were slots available in the m hospitals.

The interest, naturally, was in finding a way of assigning each student to at most one hospital, in such a way that all available positions in all hospitals were filled. (Since we are assuming a surplus of students, there would be some students who do not get assigned to any hospital.)

We say that an assignment of students to hospitals is stable if neither of the following situations arises.

- First type of instability: There are students s and s', and a hospital h, so that
    - s is assigned to h, and
    - s' is assigned to no hospital, and
    - h prefers s' to s.

- Second type of instability: There are students s and s, and hospitals h and h, so that
    - s is assigned to h, and
    - s' is assigned to h', and
    - h prefers s' to s, and
    - s' prefers h to h'.

So we basically have the Stable Matching Problem, except that (i)hospitals generally want more than one resident, and (ii) there is a surplus of medical students.

Show that there is always a stable assignment of students to hospitals,and give an algorithm to find one.

Answer

The algorithm is as following aaaaaaafasdfa:

Algorithm 1 Stable assignment algorithm

1: Initialize S to empty assignment
2: **while** There are hospitals have available positions for hiring residents **do**
3:      Pick one hospital h_i
4:      h_i invites top ranked residents r_j on the ranking list who has not been invites before
5:      **if** r_j has not been assigned **then**
6:          add pair $h_i- = r_j$ to matching S
7:      **else if** r_j prefer h_i to its current assigned hospital h_k **then**
8:          Remove pair $h_k - r_j$ from matching S
9:          Add pair $h_i- = r_j$ to matching S
10:      **else**
11:          Invite next resident
12:      **end if**
13: **end while**
14: **return** Stable assignment S

For the first type of instability, if h prefers s' to s, then, before s is assigned to h, h must have invite s'. So s' can not been assigned to no hospital.

For the second type of instability, If h prefers s' to s, then, before s is assigned to h, h must have invite s'. So the assigned hospital of s' must be higher than h on the ranking list of s'. s' can not prefer h to h'.

In conclusion, The Algorithm above can find the stable assignment.

p22, Ex1.5

The Stable Matching Problem, as discussed in the text, assumes that all men and women have a fully ordered list of preferences. In this problem we will consider a version of the problem in which men and women can be indifferent between certain options. As before we have a set M of n men and a set W of n women. Assume each man and each woman ranks the members of the opposite gender, but now we allow ties in the ranking. For example (with n = 4), a woman could say that m_1 is ranked in first place; second place is a tie between m_2 and m_3 (she has no preference between them); and m_4 is in

last place. We will say that w prefers m to m' if m is ranked higher than m' on her preference list (they are not tied).

With indifference in the rankings, there could be two natural notions for stability. And for each, we can ask about the existence of stable matching, as follows.

(a) A strong instability in a perfect matching S consists of a man m and a woman w, such that each of m and w prefers the other to their partner in S. Does there always exist a perfect matching with no strong instability? Either give an example of a set of men and women with preference lists for which every perfect matching has a strong instability; or give an algorithm that is guaranteed to find a perfect matching with no strong instability.

(b) A weak instability in a perfect matching S consists of a man m and a woman w, such that their partners in S are w and m, respectively, and one of the following holds:
- m prefers w to w, and w either prefers m to m or is indifferent between these two choices; or
- w prefers m to m, and m either prefers w to w or is indifferent between these two choices.

In other words, the pairing between m and w is either preferred by both, or preferred by one while the other is indifferent. Does there always exist a perfect matching with no weak instability? Either give an example of a set of men and women with preference lists for which every perfect matching has a weak instability; or give an algorithm that is guaranteed to find a perfect matching with no weak instability.

Answer

(a)

True. The original algorithm is still work. If m prefer w than his assigned partner a, that means m have already proposed to w before m proposed to a. If m have already proposed to w, then the assigned partner of w must be better than m because each change w get a better partner.

(b)

False. Example is show blew:

Name	Most like	Least like
A	X	Y
B	X	Y
X	AB	
Y	AB	

Wether Y married A or B, another prefer Y than X and Y is indifferent between A and B. So, there is not such an algorithm guarantee no weak instability

Homework 2

p68 Ex2.6

Consider the following basic problem. You're given an array A consisting of n integers $A[1], A[2], ..., A[n]$. You'd like to output a two-dimensional n-by-n array B in which $B[i,j](for \quad i<j)$ contains the sum of array entries $A[i]$ through $A[j]$—that is, the sum $A[i] + A[i+1] + ... + A[j]$. (The value of array entry $B[i,j]$ is left unspecified whenever $i \geq j$, so it doesn't matter what is output for these values.)

Here's a simple algorithm to solve this problem.

For i = 1, 2,..., n
    For j = i+1, i+2,..., n
        Add up array entries $A[i]$ through $A[j]$
        Store the result in $B[i,j]$
    Endfor
  Endfor

(a) For some function f that you should choose, give a bound of the form $O(f(n))$ on the running time of this algorithm on an input of size n (i.e., a bound on the number of operations performed by the algorithm).

(b) For this same function f, show that the running time of the algorithm on an input of size n is also $\Omega(f(n))$. (This shows an asymptotically tight bound of $\Theta(f(n))$ on the running time.)

(c) Although the algorithm you analyzed in parts (a) and (b) is the most natural way to solve the problem—after all, it just iterates through the relevant entries of the array B, filling in a value for each—it contains some highly unnecessary sources of inefficiency. Give a different algorithm to solve this problem, with an asymptotically better running time. In other words, you should design an algorithm with running time O(g(n)), where $\lim_{n\to\infty} \frac{g(n)}{f(n)} = 0$.

Answer

(a)

The operation of summation takes no more than n additions. So the upper bound on the number of addition operation performed by the algorithm is no more than $O(n^3)$

(b)

For i from 1 to $\frac{1}{4}n$ and j from $\frac{3}{4}n$ to n, for every $B[i,j]$, it takes at least $\frac{1}{2}n$ additions. Therefor, it takes at least $\frac{1}{4}n * \frac{1}{4}n * \frac{1}{2}n = \frac{1}{32}n^3$ additions. So, the running time of the algorithm on an input of size n is also $\Omega(n^3)$

(c)

The algorithm is as following:

Algorithm 2 Summation matrix algorithm

1: Initialize B to all zero matrix
2: **for** $i = 1, 2, ..., n$ **do**
3:     $B[i, i] = A[i]$
4:     **for** $j = i + 1, i + 2, ..., n$ **do**
5:         $B[i, j] = B[i, j - 1] + A[j]$
6:     **end for**
7: **end for**
8: **return** B

Each B[i, j] can be calculated by adding once based on B[i, j-1]. So, this algorithm has running time $O(n^2)$

p69, Ex2.8

You're doing some stress-testing on various models of glass jars to determine the height from which they can be dropped and still not break. The setup for this experiment, on a particular type of jar, is as follows. You have a ladder with n rungs, and you want to find the highest rung from which you can drop a copy of the jar and not have it break. We call this the highest safe rung.

It might be natural to try binary search: drop a jar from the middle rung, see if it breaks, and then recursively try from rung $n/4$ or $3n/4$ depending on the outcome. But this has the drawback that you could break a lot of jars in finding the answer.

If your primary goal were to conserve jars, on the other hand, you could try the following strategy. Start by dropping a jar from the first rung, then the second rung, and so forth, climbing one higher each time until the jar breaks. In this way, you only need a single jar—at the momentit breaks, you have the correct answer—but you may have to drop it n times (rather than $logn$ as in the binary search solution).

So here is the trade-off: it seems you can perform fewer drops if you're willing to break more jars. To understand better how this tradeoff works at a quantitative level, let's consider how to run this experiment given a fixed "budget" of $k \geq 1$ jars. In other words, you have to determine the correct answer—the highest safe rung—and can use at most k jars in doing so.

(a) Suppose you are given a budget of k = 2 jars. Describe a strategy for finding the highest safe rung that requires you to drop a jar at most $f(n)$ times, for some

function $f(n)$ that grows slower than linearly. (In other words, it should be the case that $lim_{n\to\infty} f(n)/n = 0$.)

(b) Now suppose you have a budget of k > 2 jars, for some given k. Describe a strategy for finding the highest safe rung using at most k jars. If $f_k(n)$ denotes the number of times you need to drop a jar according to your strategy, then the functions $f_1, f_2, f_3, ...$ should have the property that each grows asymptotically slower than the previous one: $lim_{n\to\infty} f_k(n)/f_{k-1}(n) = 0$ for each k.

Answer

(a)

The algorithm is as following:

Algorithm 3 Algorithm of Ex2.8(a)

1: Divide n rungs into $\sqrt{n}$ parts. Each part has $\sqrt{n}$ rungs.
2: Start by dropping the first jar from the highest rung of the first rung part.
3: **repeat**
4:    Climbing $\sqrt{n}$ rungs higher to the highest rung of the next part.
5: **until**
6:    The glass jar breaks
7: Start by dropping the second jar from the first rung of the part in which the first jar has broken.
8: **repeat**
9:    Climbing one rung higher each time
10: **until**
11:    The glass jar breaks
12: **return** The height

It takes at most $2\sqrt{n}$ times test to determine the height. So, this algorithm has running time $O(\sqrt{n})$

(b)

The algorithm is as following:

For a given k, it takes $O(kn^{\frac{1}{k}})$ times test to determine the height. So,

$$lim_{n\to\infty} f_k(n)/f_{k-1}(n) = lim_{n\to\infty} \frac{kn^{\frac{1}{k}}}{(k-1)n^{\frac{1}{k-1}}} = 0$$

Algorithm 4 Algorithm of Ex2.8(b)

1: Divide n rungs into $n^{\frac{1}{k}}$ parts. Each part has $n^{\frac{k-1}{k}}$ rungs. Then, divide every part into $n^{\frac{1}{k}}$ smaller parts and so on. We totally have

2: **for** $i = 1, 2, ..., k-1$ **do**

3:    Start by dropping the i^{th} jar from the highest rung of the first rung part in this level of partition.

4:    **repeat**

5:     Climbing $n^{\frac{k-i}{k}}$ rungs higher to the highest rung of the next part in this level of partition.

6:    **until**

7:     The glass jar breaks

8:    Do tests in this part in which last jar has broken by next level of partition.

9: **end for**

10: Start by dropping the k^{th} jar from the first rung of the part of this level, in which last jar has broken.

11: **repeat**

12:    Climbing one rung higher each time

13: **until**

14:    The glass jar breaks

15: **return** The height
---

p107 Ex3.5

A binary tree is a rooted tree in which each node has at most two children. Show by induction that in any binary tree the number of nodes with two children is exactly one less than the number of leaves.

Answer

Suppose T is a binary tree with n_1 nodes with one child, n_2 nodes with two children and m leaves without child. There are two ways to add a node.

- Add a leaf to a leaf node. Then, the new tree will have $n_1 + 1$ nodes with one child n_2 nodes with two children and $m - 1 + 1 = m$ leaves without child.

- Add a leaf to a node with one child. Then, the new tree will have $n_1 - 1$ nodes with one child $n_2 + 1$ nodes with two children and $m + 1$ leaves without child.

That means adding node will increase the number of n_2 and m simultaneously. For the simplest tree which only have one root node, $n_2 = 0$, $m = 1$, And every binary tree can be built from root node by adding nodes one by one. So, in any binary tree the number of nodes with two children is exactly one less than the number of leaves.

p107 Ex3.6

We have a connected graph $G = (V, E)$, and a specific vertex $u \in V$. Suppose we compute a depth-first search tree rooted at u, and obtain a tree T that includes all nodes of G. Suppose we then compute a breadth-first search tree rooted at u, and obtain the same tree T. Prove that $G = T$. (In other words, if T is both a depth-first search tree and a breadth-first search tree rooted at u, then G cannot contain any edges that do not belong to T.)

Answer

By contradiction, suppose there exists an edge $e = \{A - B\} \in G$ but $e = \{A - B\} \notin T$. As T is a tree, so there must be a path $\{A - C - B\}$ from A to B. C is a vertex different from A and B. According to the distance between A, B, C and the root u in T, there are two conditions:

- $d_T(A, u), d_T(B, u) > d_T(C, u)$. Then, C is the parent node of both A and B. In the process of constructing DFS tree, after C has been added to T, no matter A or B is added to T first, $\{A - B\} \in T$.

- $d_T(A, u) > d_T(C, u) > d_T(B, u)$ (or $d_T(B, u) > d_T(C, u) > d_T(A, u)$). Then, B is the parent node of both A and C. In the process of constructing BFS tree, after B has been added to T, as C haven't been added to T, A will be added to T after B. then, $\{A - B\} \in T$

In conclusion, $G = T$.

10

Homework 3

p107 Ex3.2

The algorithm described in Section 3.6 for computing a topological ordering of a DAG repeatedly finds a node with no incoming edges and deletes it. This will eventually produce a topological ordering, provided that the input graph really is a DAG.

But suppose that we're given an arbitrary graph that may or may not be a DAG. Extend the topological ordering algorithm so that, given an input directed graph G, it outputs one of two things: (a) a topological ordering, thus establishing that G is a DAG; or (b) a cycle in G, thus establishing that G is not a DAG. The running time of your algorithm should be $O(m+n)$ for a directed graph with n nodes and m edges.

Answer

The Algorithm is on the next page

Topological ordering is $O(m+n)$ and DFS search is $O(m+n)$. So, the extend algorithm is $O(m+n)$

p109 Ex3.8

A number of stories in the press about the structure of the Internet and the Web have focused on some version of the following question: How far apart are typical nodes in these networks? If you read these stories carefully, you find that many of them are confused about the difference between the diameter of a network and the average distance in a network; they often jump back and forth between these concepts as though they're the same thing.

As in the text, we say that the distance between two nodes u and v in a graph G = (V, E) is the minimum number of edges in a path joining them; we'll denote this by dist(u, v). We say that the diameter of G is the maximum distance between any pair of nodes; and we'll denote this quantity by diam(G).

Let's define a related quantity, which we'll call the average pairwise distance in G (denoted apd(G)). We define apd(G) to be the average, over all $\binom{n}{2}$ sets of two distinct nodes u and v, of the distance between u and v. That is,

$$\text{apd}(G) = \left[\sum_{\{u,v\} \subseteq V} \text{dist}(u,v) \right] / \binom{n}{2}$$

while

$$\text{apd}(G) = [\text{dist}(u,v) + \text{dist}(u,w) + \text{dist}(v,w)]/3 = 4/3$$

Algorithm 5 Algorithm of Ex3.2

1: Initialize empty ordering O and empty circle C

2: **while** There is a node $v \in G$ with no incoming edge **do**

3:     Add the node v into O.

4:     Delete v and connected edges from G and get new $G = G - \{v\}$

5: **end while**

6: **if** $G == \emptyset$ **then**

7: **return** The topological ordering O

8: **else**

9:     Mark all nodes as unreached

10:     Apply DFS Search from any node v

11:     **function** DFS(G,v)

12:         Mark v as being reaching

13:         **if** There is a path from v to an being reaching node $u \in G$ **then**

14:             There is a circle $u \to v \to u$

15:             Add u to the end of circle C

16:             Return{Circle is found}

17:         **end if**

18:         **while** There is a path from v to an unreached node $u \in G$ **do**

19:             Apply DFS to node u: (DFS(G,u))

20:             **if** Return "Circle is found" **then**

21:                 Add u to the end of circle C

22:                 Return{Circle is found}

23:             **end if**

24:         **end while**

25:         Mark v as reached

26:         Return

27:     **end function**

28: **end if**

29: **return** The Circle C

Of course, these two numbers aren't all that far apart in the case of this three-node graph, and so it's natural to ask whether there's always a close relation between them. Here's a claim that tries to make this precise.

Claim: There exists a positive natural number c so that for all connected graphs G, it is the case that

$$\frac{\operatorname{diam}(G)}{\operatorname{apd}(G)} \leq c$$

Decide whether you think the claim is true or false, and give a proof of either the claim or its negation.

Answer

False. Counterexample is described below:

There is a path in G from v_1 to $v_{\sqrt{n}}$ connected $\sqrt{n}$ nodes; Other $n - \sqrt{n}$ nodes are connected to v_1. $diam(G)$ is $O(\sqrt{n})$, $\sum_{\{u,v\}\subseteq V} \text{dist}(u,v)$ is $O(n^2)$, $\binom{n}{2}$ is $O(n^2)$, so, $apd(G)$ is $O(1)$. As $n \to \infty$, $\frac{\text{diam}(G)}{\text{apd}(G)} \to \infty$

p190 Ex4.4

Some of your friends have gotten into the burgeoning field of time-series data mining, in which one looks for patterns in sequences of events that occur over time. Purchases at stock exchanges—what's being bought—are one source of data with a natural ordering in time. Given a long sequence S of such events, your friends want an efficient way to detect certain "patterns" in them—for example, they may want to know if the four events

    buy Yahoo, buy eBay, buy Yahoo, buy Oracle

occur in this sequence S, in order but not necessarily consecutively.

They begin with a collection of possible events (e.g., the possible transactions) and a sequence S of n of these events. A given event may occur multiple times in S (e.g., Yahoo stock may be bought many times in a single sequence S). We will say that a sequence S' is a subsequence of S if there is a way to delete certain of the events from S so that the remaining events, in order, are equal to the sequence S'. So, for example, the sequence of four events above is a subsequence of the sequence

    buy Amazon, buy Yahoo, buy eBay, buy Yahoo, buy Yahoo, buy Oracle

Their goal is to be able to dream up short sequences and quickly detect whether they are subsequences of S. So this is the problem they pose to you: Give an algorithm that takes two sequences of events—S' of length m and S of length n, each possibly containing an event more than once—and decides in time $O(m+n)$ whether S' is a subsequence of S.

Answer

The Algorithm is as following:

We can know the result after the pointer traversing the entire sequence. So, the algorithm is $O(m+n)$.

Algorithm 6 Algorithm of Ex4.4

1: Initialize pointer at location 0 of S
2: **for** $i \in S'$ **do**
3:     **if** next events j appear in S after pointer **then**
4:         Change pointer to the location of j
5:     **else**
6:         Return False
7:     **end if**
8: **end for**
9: **return** True

p191 Ex4.6

Your friend is working as a camp counselor, and he is in charge of organizing activities for a set of junior-high-school-age campers. One of his plans is the following mini-triathalon exercise: each contestant must swim 20 laps of a pool, then bike 10 miles, then run 3 miles. The plan is to send the contestants out in a staggered fashion, via the following rule: the contestants must use the pool one at a time. In other words, first one contestant swims the 20 laps, gets out, and starts biking. As soon as this first person is out of the pool, a second contestant begins swimming the 20 laps; as soon as he or she is out and starts biking, a third contestant begins swimming . . . and so on.) Each contestant has a projected swimming time (the expected time it will take him or her to complete the 20 laps), a projected biking time (the expected time it will take him or her to complete the 10 miles of bicycling), and a projected running time (the time it will take him or her to complete the 3 miles of running). Your friend wants to decide on a schedule for the triathalon: an order in which to sequence the starts of the contestants. Let's say that the completion time of a schedule is the earliest time at which all contestants will be finished with all three legs of the triathalon, assuming they each spend exactly their projected swimming, biking, and running times on the three parts. (Again, note that participants can bike and run simultaneously, but at most one person can be in the pool at any time.) What's the best order for sending people out, if one wants the whole competition to be over as early as possible? More precisely, give an efficient algorithm that produces a schedule whose completion time is as small as possible.

Answer

The strategy is letting contestant who takes longer time for biking and 3 miles begin to swim earlier in the schedule sequence. We can approve this is the optimal algorithm by using contradiction:

Suppose A takes longer time for biking and 3 miles than B, but is after B in the schedule sequence. Without loss of generality, let's assume A is next to B. Suppoese A and B take $t1_A, t1_B, t2_A, t2_B, t3_A, t3_B$ for swimming, biking and 3 miles. It takes $t1_B + t1_A + t2_A + t3_A$ for them

to finish. However, if We let A swim before B, it takes $t1_A + maximum\{(t2_A + t3_A), (t1_B + t2_B + t3_B)\}$

$$\begin{aligned} t_{sum} &= t1_A + maximum\{(t2_A + t3_A), (t1_B + t2_B + t3_B)\} \\ &\leq t1_A + maximum\{(t2_A + t3_A), (t1_B + t2_A + t3_A)\} \\ &\leq t1_B + t1_A + t2_A + t3_A \end{aligned} \tag{1}$$

So, the algorithm is optimal.

Homework 4

p192 Ex4.8

Suppose you are given a connected graph G, with edge costs that are all distinct. Prove that G has a unique minimum spanning tree.

Answer

By contradiction, suppose T and T' are two distinct minimum spanning trees of G. There must be an edge $e \in T', e \notin T$, If we add e to T, a circle will be formed. If e is not the most expensive edge, it is condradicting the fact that T is a MST. If e is the most expensive edge, there exists an edge $e' \in T, e' \notin T'$ in the circle. If we add e' to T', a circle contain e and e' will be formed. it is condradicting the fact that T' is a MST.

p197 Ex4.18

Your friends are planning an expedition to a small town deep in the Canadian north next winter break. They've researched all the travel options and have drawn up a directed graph whose nodes represent intermediate destinations and edges represent the roads between them.

In the course of this, they've also learned that extreme weather causes roads in this part of the world to become quite slow in the winter and may cause large travel delays. They've found an excellent travel Web site that can accurately predict how fast they'll be able to travel along the roads; however, the speed of travel depends on the time of year. More precisely, the Web site answers queries of the following form: given an edge $e = (v, w)$ connecting two sites v and w, and given a proposed starting time t from location v, the site will return a value $f_e(t)$, the predicted arrival time at w. The Web site guarantees that $f_e(t) \geq t$ for all edges e and all times t (you can't travel backward in time), and that $f_e(t)$ is a monotone increasing function of t (that is, you do not arrive earlier by starting later). Other than that, the functions $f_e(t)$ may be arbitrary. For example, in areas where the travel time does not vary with the season, we would have $f_e(t) = t + l_e$, where l_e is the time needed to travel from the beginning to the end of edge e.

Your friends want to use the Web site to determine the fastest way to travel through the directed graph from their starting point to their intended destination. (You should assume that they start at time 0, and that all predictions made by the Web site are completely correct.) Give a polynomial-time algorithm to do this, where we treat a single query to the Web site (based on a specific edge e and a time t) as taking a single computational step.

Answer

The Algorithm is as following:

Algorithm 7 Algorithm of Ex4.18

1: Initialize explored nodes set $S = s$, earliest arrival time of s $d(s) = 0$ and last node of s $r(s) = \emptyset$
2: **while** $S \neq V$ **do**
3:     Compute the distances from nodes inside S to nodes outside S and have a direct edge. Each distance d_{ab} is decided by a, b and the distace from a to s
4:     Find the shortest distance d_{ab}, add b to S and define $d(b) = d_a + d_{ab}$ $r(b) = a$
5: **end while**
6: **return**

The Correctness is guaranteed because that $f_e(t)$ is a monotone increasing function of t. For any path from a to b, earlier reach to a, earlier reach to b. The algorithm above computer the earliest arrival time one node by one node.

Similar to DijKstra's algorithm. This algorithm is $O(mlogn)$

p246 Ex5.1

You are interested in analyzing some hard-to-obtain data from two separate databases. Each database contains n numerical values—so there are 2n values total—and you may assume that no two values are the same. You'd like to determine the median of this set of 2n values, which we will define here to be the n_{th} smallest value.

However, the only way you can access these values is through queries to the databases. In a single query, you can specify a value k to one of the two databases, and the chosen database will return the kth smallest value that it contains. Since queries are expensive, you would like to compute the median using as few queries as possible. Give an algorithm that finds the median value using at most $O(logn)$ queries.

Answer

First, we can compare the median of the two database. If $A[\lfloor \frac{n}{2} \rfloor] < B[\lfloor \frac{n}{2} \rfloor]$, we can know the median number is on higher position of A and lower position of B. So we can reduce half number of the two databse.

The Algorithm is as following, evaluate MEDIAN(A[n],B[n],n):

For each iteration, each database is queried once. It will iterate for a maximum of $\lceil log_2 n \rceil$ times. So, the algorithm is $O(logn)$.

Algorithm 8 Algorithm of Ex5.1

1: **function** MEDIAN(a[n],b[n],k)
2:    **if** k==1 **then**
3:        Return $min\{a[1], b[1]\}$
4:    **else if** $a[\lfloor\frac{k}{2}\rfloor] < b[\lfloor\frac{k}{2}\rfloor]$ **then**
5:        Return MEDIAN(a[$\lfloor\frac{k}{2}\rfloor$:n], b[1:$\lfloor\frac{k}{2}\rfloor$],$\lceil\frac{k}{2}\rceil$)
6:    **else**
7:        Return MEDIAN(a[1:$\lfloor\frac{k}{2}\rfloor$], b[$\lfloor\frac{k}{2}\rfloor$:n],$\lceil\frac{k}{2}\rceil$)
8:    **end if**
9: **end function**

p246 Ex5.2

Recall the problem of finding the number of inversions. As in the text, we are given a sequence of n numbers $a_1, ..., a_n$, which we assume are all distinct, and we define an inversion to be a pair $i < j$ such that $a_i > a_j$.

We motivated the problem of counting inversions as a good measure of how different two orderings are. However, one might feel that this measure is too sensitive. Let's call a pair a significant inversion if $i < j$ and $a_i > 2a_j$. Give an $O(nlogn)$ algorithm to count the number of significant inversions between two orderings.

Answer

The Algorithm is as following, evaluate INVERSION(a[n]):

Algorithm 9 Algorithm of Ex5.2

1: **function** INVERSION(a[n])
2:    **if** n==1 **then** Return 0
3:    **else**
4:        x = $a[1 : \lfloor\frac{n}{2}\rfloor]$; y = $a[\lceil\frac{n}{2}\rceil : n]$
5:        Initialize pointers i and j at the first index of x and y; counter=0;
6:        **while** i,j are not at end of x and y **do**
7:            **if** $a_i > 2a_j$ **then**
8:                counter = counter + $\lfloor\frac{n}{2}\rfloor - i + 1$; j++
9:            **else**
10:                i++
11:            **end if**
12:        **end while**
13:        Return INVERSION(x)+INVERSION(y)+counter
14:    **end if**
15: **end function**

The algorithm is $O(nlogn)$.

Homework 5

Prove the Master Theorem

Suppose that $T(n)$ is a function on the nonnegative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + f(n)$$

where $\frac{n}{b}$ means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Let $k = \log_b a$,

- Case 1. If $f(n) = O\left(n^{k-\epsilon}\right)$ for some constant $\epsilon > 0$, then $T(n) = \Theta\left(n^k\right)$
- Case 2. If $f(n) = \Theta\left(n^k \log^p n\right)$, then $T(n) = \Theta\left(n^k \log^{p+1} n\right)$
- Case 3. If $f(n) = \Omega\left(n^{k+\epsilon}\right)$ for some constant $\epsilon > 0$ and if $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n, then $T(n) = \Theta(f(n))$

Answer

By mathematical induction, all the three cases are valid for $n = 2$. Suppose they are valid for $n = m$.

- Case 1. If $f(n) = O\left(n^{k-\epsilon}\right)$, $T(n) = \Theta\left(n^k\right)$,

$$T(bm) = aT(m) + f(bm) = b^k\Theta\left(m^k\right) + O\left((bm)^{k-\epsilon}\right) = \Theta\left((bm)^k\right)$$

- Case 2. If $f(n) = \Theta\left(n^k \log^p n\right)$, $T(n) = \Theta\left(n^k \log^{p+1} n\right)$,

$$T(bm) = b^k\Theta\left(m^k \log^{p+1} m\right) + \Theta\left((bm)^k \log^p bm\right)$$

$$= \Theta\left((bm)^k(log^{p+1}m + \log^p bm)\right) = \Theta\left((bm)^k \log^{p+1} bm\right)$$

- If $f(n) = \Omega\left(n^{k+\epsilon}\right)$ for some constant $\epsilon > 0$, $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n and if $T(n) = \Theta(f(n))$,

$$T(bm) = b^k\Omega\left(m^{k+\epsilon}\right) + \Theta(\Omega\left((bm)^{k+\epsilon}\right)) \leq c\Omega\left(m^{k+\epsilon}\right) + \Theta(\Omega\left((bm)^{k+\epsilon}\right))$$

so, $T(bm) = \Theta(f(bm))$

So far, we have already proved the condition of $n = b^k$. For any other integer, we can always extend it to $n = b^k$ without affecting the time complexity.

p246 Ex5.3

Suppose you're consulting for a bank that's concerned about fraud detection, and they come to you with the following problem. They have a collection of n bank cards that they've confiscated, suspecting them of being used in fraud. Each bank card is

a small plastic object, containing a magnetic stripe with some encrypted data, and it corresponds to a unique account in the bank. Each account can have many bank cardscorresponding to it, and we'll say that two bank cards are equivalent if they correspond to the same account.

It's very difficult to read the account number off a bank card directly, but the bank has a high-tech "equivalence tester" that takes two bank cards and, after performing some computations, determines whether they are equivalent.

Their question is the following: among the collection of n cards, is there a set of more than $n/2$ of them that are all equivalent to one another? Assume that the only feasible operations you can do with the cards are to pick two of them and plug them in to the equivalence tester. Show how to decide the answer to their question with only $O(n \log n)$ invocations of the equivalence tester.

Answer

The Algorithm is as following:

Algorithm 10 Algorithm of Ex5.3

1: **function** F(n cards)
2:     **if** n==1 **then**
3:         ReturnThis card
4:     **else if** F(Half of these n cards) or F(The other half of these n cards) $\neq$ NONE **then**
5:         Test any of the returned m casds with all the cards in the other half.
6:         **if** There are more than $n/2 - m$ equivalent cards in the other half **then**
7:             Return m' equivalent cards
8:         **end if**
9:     **else**
10:         Return NONE
11:     **end if**
12: **end function**

The algorithm is $O(logn)$.

p247 Ex5.4

You've been working with some physicists who need to study, as part of their experimental design, the interactions among large numbers of very small charged particles. Basically, their setup works as follows. They have an inert lattice structure, and they use this for placing charged particles at regular spacing along a straight line. Thus we can model their structure as consisting of the points $\{1, 2, 3, \ldots, n\}$ on the real line;

and at each of these points j, they have a particle with charge q_j. (Each charge can be either positive or negative.)

They want to study the total force on each particle, by measuring it and then comparing it to a computational prediction. This computational part is where they need your help. The total net force on particle j, by Coulomb's Law, is equal to

$$F_j = \sum_{i<j} \frac{Cq_iq_j}{(j-i)^2} - \sum_{i>j} \frac{Cq_iq_j}{(j-i)^2}$$

They've written the following simple program to compute F_j for all j

```
For j = 1, 2, ... , n
   Initialize F_j to 0
   For i = 1, 2, ..., n
      If i < j then
         Add (C q_i q_j)/((j - i)^2) to F_j
      Else if i > j then
         Add -(C q_i q_j)/((j - i)^2) to F_j
      Endif
   Endfor
   Output F_j
Endfor
```

The trouble is, for the large values of n they're working with, the program takes several minutes to run. On the other hand, their experimental setup is optimized so that they can throw down n particles, perform the measurements, and be ready to handle n more particles within a few seconds. So they'd really like it if there were a way to compute all the forces F_j much more quickly, so as to keep up with the rate of the experiment. Help them out by designing an algorithm that computes all the forces F_j in $O(n \log n)$ time.

Answer

We can first construct a vector $\vec{a} = \left(\frac{1}{n^2}, \ldots, \frac{1}{4}, 1, 0, -1, -\frac{1}{4}, \ldots - \frac{1}{n^2}\right)$. It takes $O(n)$ time. Assume $\vec{q} = (q_1, q_2, \ldots, q_n)$. Then we can use FFT to compute the convolution of $\vec{a} \otimes \vec{q}$. The total time is $O(nlogn)$.

Homework 6

p248 Ex5.5

Hidden surface removal is a problem in computer graphics that scarcely needs an introduction: when Woody is standing in front of Buzz, you should be able to see Woody but not Buzz; when Buzz is standing in front of Woody, . . . well, you get the idea.

The magic of hidden surface removal is that you can often compute things faster than your intuition suggests. Here's a clean geometric example to illustrate a basic speed-up that can be achieved. You are given n nonvertical lines in the plane, labeled $L_1, \ldots, L_n$, with the t^{h} line specified by the equation $y = a_i x + b_i$. We will make the assumption that no three of the lines all meet at a single point. We say line L_i is uppermost at a given x -coordinate x_0 if its y -coordinate at x_0 is greater than the y -coordinates of all the other lines at $x_0 : a_i x_0 + b_i > a_j x_0 + b_j$ for all $j \neq i$. We say line L_i is visible if there is some x -coordinate at which it is uppermost-intuitively, some portion of it can be seen if you look down from " $y = \infty$."

Give an algorithm that takes n lines as input and in $O(n \log n)$ time returns all of the ones that are visible. Figure 5.10 gives an example.

Answer

The Algorithm is as following:

Algorithm 11 Algorithm of Ex5.5

1: Sort all lines by slope
2: **function** F(n ordered lines)
3:    **if** n==2 **then**
4:        Return The intersection and two line
5:    **else**
6:        L1 = F(The first half of these n lines)
7:        L2 = F(The other half of these n lines)
8:        Slope of lines in L1 is always smaller than slope of lines in L2
9:        L1 and L2 are two broken line with p-1 intersection points, p segments and q-1 intersection points, q segments. As the x coordinate of these p-1 and q-1 intersection points are in order, it takes no more than O(p+q) time to find the intersection point k of these two broken line.
10:        After we have found k, lines(segments) in L1 with smaller endpoint(intersection point) x coordinate are visible; lines(segments) in L2 with larger endpoint(intersection point) x coordinate are visible.
11:        Return Visible lines and intersection points in order
12:    **end if**
13: **end function**

$p + q$ is always smaller than n, so the Algorithm is $O(nlogn)$

p248 Ex5.6

Consider an n-node complete binary tree T, where $n = 2^d - 1$ for some d. Each node v of T is labeled with a real number x_v. You may assume that the real numbers labeling the nodes are all distinct. A node v of T is a local minimum if the label x_v is less than the label x_w for all nodes w that are joined to v by an edge.

You are given such a complete binary tree T, but the labeling is only specified in the following implicit way: for each node v, you can determine the value x_v by probing the node v. Show how to find a local minimum of T using only $O(\log n)$ probes to the nodes of T.

Answer

We probe node from root. If the two son is bigger, than the root is the local minimum. Otherwise, if the one of the two son is smaller, we choose any smaller son and repeat the above. As each time we choose a son node smaller than the parent, we stop at a local minimum, or we stop at a leaf node. The leaf node is also satisfied the definition of local minimum. The algorithm is abviously $O(logn)$.

p317 Ex6.6

Answer

The Algorithm is as following:

$p + q$ is always smaller than n, so the Algorithm is $O(nlogn)$

p319 Ex6.8

Answer

(a)

False. Example is show blew:

i	1	2	3	4
x_i	1	10	10	2
$f(i)$	1	2	4	8

The algorithm in (a) will return 4. However, the correct answer is a total of 5.

Algorithm 12 Algorithm of Ex6.6

1: **function** F(first n words)
2:    **if** n==0 **then**
3:       Return 0
4:    **else**
5:       **for** i = 1 to n **do**
6:          F(i) = ∞
7:          **for** j = 1 to i **do**
8:             **if** $\left[\sum_{i=j}^{n-1} (c_i + 1)\right] + c_n \leq L$ **then**
9:                $slack(j, n) = L - \left[\sum_{i=j}^{n-1} (c_i + 1)\right] - c_n$
10:            **else**
11:               slack(j,n) = ∞
12:            **end if**
13:            **if** $slack(j, n)$+F(first $j - 1$ words) < F(i) **then**
14:               F(i) = $slack(j, n)$+F(first $j - 1$ words)
15:            **end if**
16:         **end for**
17:      **end for**
18:   **end if**
19: **return** F(n)
20: **end function**

(b)

The Algorithm is as following:

Algorithm 13 Algorithm of Ex6.8

1: **function** F(n)
2:    **if** n==0 **then**
3:       Return 0
4:    **else**
5:       **for** i = 1 to n **do**
6:          F(i) = 0
7:          **for** j = 0 to i **do**
8:             s(j,i) = $min\{\min(x_i, f(j))\}$
9:             **if** $s(j, i)$+F(j) > F(i) **then**
10:               F(i) = $s(j, i)$+F(j)
11:            **end if**
12:         **end for**
13:      **end for**
14:   **end if**
15: **end function**

Homework 7

p328 Ex6.18

Consider the sequence alignment problem over a four-letter alphabet (z_1, z_2, z_3, z_4), with a given gap cost and given mismatch costs. Assume that each of these parameters is a positive integer.

Suppose you are given two strings $A = a_1 a_2 \cdots a_m$ and $B = b_1 b_2 \cdots b_n$ and a proposed alignment between them. Give an $O(mn)$ algorithm to decide whether this alignment is the unique minimum-cost alignment between A and B.

Answer

In the alignment algorithm

$$
OPT(i, j) = \begin{cases} j\delta & \text{if } i = 0 \\ \min \begin{cases} \alpha_{x_i y_j} + OPT(i-1, j-1) \\ \delta + OPT(i-1, j) \\ \delta + OPT(i, j-1) \end{cases} & \text{otherwise} \\ i\delta & \text{if } j = 0 \end{cases}
$$

We can check $\alpha_{x_i y_j} + OPT(i-1, j-1), \quad \delta + OPT(i-1, j), \quad \delta + OPT(i, j-1)$ each time. If in every iteration, the minimum is unique, then the alignment is unique.

p329 Ex6.20

Suppose it's nearing the end of the semester and you're taking n courses, each with a final project that still has to be done. Each project will be graded on the following scale: It will be assigned an integer number on a scale of 1 to $g > 1$, higher numbers being better grades. Your goal, of course, is to maximize your average grade on the n projects.

You have a total of $H > n$ hours in which to work on the n projects cumulatively, and you want to decide how to divide up this time. For simplicity, assume H is a positive integer, and you'll spend an integer number of hours on each project. To figure out how best to divide up your time, you've come up with a set of functions $\{f_t : i = 1, 2, \ldots, n\}$ (rough estimates, of course) for each of your n courses; if you spend $h \leq H$ hours on the project for course i, you'll get a grade of $f_t(h)$. (You may assume that the functions f_i are nondecreasing: if $h < h'$, then $f_t(h) \leq f_t(h')$.)

So the problem is: Given these functions $\{f_t\}$, decide how many hours to spend on each project (in integer values only) so that your average grade, as computed according to the f_b, is as large as possible. In order to be efficient, the running time of your algorithm should be polynomial in n, g, and H; none of these quantities should appear as an exponent in your running time.

Answer

The Algorithm is as following:

Algorithm 14 Algorithm of Ex6.20

1: **function** $F(H, \{f_t\}, g\ )$
2:   $G(i,j) = 0, i = 1...n, j = 1...H$
3:   **for** i = 1 to n **do** G(i,1) = $Max_{j\leq n}\{f_i(1)\}$
4:     **for** (**do**j = 2 to H)
5:       G(i,j) = $Max_{k\leq j}\{G(i-1, j-k) + f_i(k)\}$
6:     **end for**
7:   **end for**
8: **return** G(n,H)
9: **end function**

---

p416 Ex7.4

Decide whether you think the following statement is true or false. If it is true, give a short explanation. If it is false, give a counterexample.

Let G be an arbitrary flow network, with a source s, a sink t, and a positive integer capacity c_e on every edge e. If f is a maximum $s - t$ flow in G, then f saturates every edge out of s with flow (i.e., for all edges e out of s, we have $f(e) = c_e$)

Answer

False, For such a flow network $\{s - v - t\}$, $c_{s-v} = 2, c_{v-t} = 1$. The maximum s-t flow doesn't saturate edge s-v.

p416 Ex7.5

Decide whether you think the following statement is true or false. If it is true, give a short explanation. If it is false, give a counterexample.

Let G be an arbitrary flow network, with a source s,| a sink t, and a positive integer capacity c_e on every edge e; and let (A, B) be a mimimum s-t cut with respect to these capacities $\{c_e : e \in E\}$. Now suppose we add 1 to every capacity; then (A, B) is still a minimum s -t cut with respect to these new capacities $\{1 + c_e : e \in E\}$

Answer

False, the counter example is as below:

$$G = \{s, v_1, v_2, v_3, v_4, t; \quad s - v_1, s - v_2, s - v_3, v_1 - v_4, v_2 - v_4, v_3 - v_4, v_4 - t\}$$

$$c_{v_4 - t} = 4, \quad otherwise, c_e = 1$$

In this example, minmum cut is (s,G-s). After we add 1 to every capacity, the minmum cut is (G-t,t)

Homework 8

p420 Ex7.12

Consider the following problem. You are given a flow network with unitcapacity edges: It consists of a directed graph $G = (V, E)$, a source $s \in V$, and a sink $t \in V$; and $c_e = 1$ for every $e \in E$. You are also given a parameter k.

The goal is to delete k edges so as to reduce the maximum $s - t$ flow in G by as much as possible. In other words, you should find a set of edges $F \subseteq E$ so that $|F| = k$ and the maximum $s \cdot t$ flow in $G' = (V, E - F)$ is as small as possible subject to this.

Give a polynomial-time algorithm to solve this problem.

Answer

After the maximum s-t flow algorithm, we can perform BFS from s in residual graph to get A which consists all nodes s can search. So we get a minimum s-t cut (A,G-A), there is at least f edges go out of A in G. We can choose any k edges of these f edges to delete:

- As these k edges are maximum flow (minimum cut) edges, so every deletion will cause a reduction of maximum s-t flow.

- As the capacity of all edges is one, every deletion can cause no more than one reduction of maximum s-t flow.

So, the algorithm is optimal. When C = 1, the running time of Ford-Fulkerson is O(mn). So the algorithm is polynomial-time.

p424 Ex7.18

Answer

(a)

Assume M'' is the largest matching. If there exists $y \in Y$ covered by M, but not covered by M''. For any node $x \in X$ has an edge to y, if x is not matched in M'', than match $x - y$ is available, that is contraydictory to M'' is the largest matching. So, there must exist a match between x and $y' \in Y$. We can always remove match $x - y'$ and add match $x - y$ without reducing number of matches.

So, for given G,M,k, We can first compute largest matching M'' of G, K_2 is the size of the M''. If K_2 is k more than K_M, than we can change M'' into M''' by the above strategy to let every node $y \in Y$ that is covered by M is also covered by M'''. After that, cut redundant matches we can get the solution M'. Otherwise, if K_2 is not k more than K_M, report there is no solution.

The running time is as same as largest matching algorithm ($O(mn)$, polynomial).

(b)

$$G = \{x_1, x_2, y_1, y_2; \quad x_1 - y_1, x_1 - y_2, x_2 - y_2\}$$

$$\{M = x_1 - y_2\}$$

When k=1, $\{M' = x_1 - y_1, x_2 - y_2\}$, the edges of M do not form a subset of the edges of M'.

(c)

By the analysis in (a), we can always change M'' into M' without reducing number of matches.

Homework 9

p507 Ex8.6

Answer

We can prove this problem X of *Monotone Satisfiability with Few True Variables* is NP-complete by proving that $Y \leq_p X$, where Y is *Vertex Cover* problem, which is NP-complete.

For any clauses consist of two variables, we can construct a graph by adding an edge $\{x_i - x_j\}$ if $\exists C_k = (x_i \vee x_j)$. Or, we can construct clauses from any graph. If there is a satisfying assignment for the instance in which at most k variables are set to 1, then we can choose these k nodes corresponding to these variables to cover all the edge.

So, problem X of *Monotone Satisfiability with Few True Variables* is NP-complete.

p517 Ex8.23

Answer

Assume string u is a concatenation with minimum length. The p^{th} position of u is equal to k^{th} position of a_i and l^{th} position of b_j. Assume L is the maximum length of strings in $A \cup B$, so, there are at most mnL^2 cases. If the length of u is longer than mnL^2, then there exists position p and p' belong to same a_i and b_j. So, we can delete all number between position p and p' and get a new concatenation u'. u' is bounded by mnL^2, where L is the maximum length of strings in $A \cup B$.

p505 Ex8.3

Answer

Similar to Ex8.6, Each sport can be regarded as a clause consists of all counselor qualified in it. $S_i = (c_{i_1} \vee c_{i_2} \vee ... \vee c_{i_k})$.

For any clauses consist of two counselors, we can construct a graph by adding an edge $\{c_i - c_j\}$ if $\exists S_k = (c_i \vee c_j)$. Or, we can construct clauses from any graph. If we can hire at most k of the counselors and have at least one counselor qualified in each of the n sports, then we can choose these k nodes corresponding to these counselors to cover all the edge. So, $Y \leq_p X$ is satisfied.

As a result, Efficient Recruiting is NP-complete.

p517 Ex8.22

Answer

We can use BFS to find all connected componenets. Then, we can call A to compute the size of independent set of each component. At last, Add them all up and get the size of independent set of graph G. As BFS is $O(|V|)$ and A runs in time polynomial. The totally running time is polynomial.

Homework 10

p507 Ex8.7

since the 3 -Dimensional Matching Problem is NP-complete, it is natural to expect that the corresponding 4 -Dimensional Matching Problem is at least as hard. Let us define 4 -Dimensional Matching as follows. Given sets W, X, Y, and Z, each of size n, and a collection C of ordered 4 -tuples of the form (w_i, x_j, y_k, z_ℓ), do there exist n 4-tuples from c so that no two have an element in common?

Prove that 4 -Dimensional Matching is NP-complete.

Answer

We can prove *4-Dimensional Matching problem* X is NP-complete by proving $Y \leq_p X$, where Y is *3-Dimensional Matching problem*, which is NP-complete.

Given disjoint sets X, Y, and Z, each of size n, and given a set $T \subseteq X \times Y \times Z$ of ordered triples, we can construct a set M disjoint from X, Y, and Z and contruct tetrad set $T' = T \times M$, $T' \subseteq X \times Y \times Z \times M$

If there exist a set of n tetrad in T' so that each element of $X \cup Y \cup Z \cup M$ is contained in exactly one of these tetrad. We can take out first three dimension of these n tetrad and get a set of n triples in T so that each element of $X \cup Y \cup Z$ is contained in exactly one of these triples?

So, $Y \leq_p X$ and *4-Dimensional Matching problem* X is NP-complete.

p518 Ex8.27

Consider the following problem. You are given positive integers $x_1, ..., x_n$, and numbers k and B. You want to know whether it is possible to partition

the numbers $\{x_i\}$ into k sets $S_1, \ldots, s_k$ so that the squared sums of the sets add up to at most B :

$$\sum_{i=1}^{k} \left(\sum_{x_j \in S_i} x_j \right)^2 \leq B$$

Show that this problem is NP-complete.

Answer

As we know, *Subset Sum Problem Y* is is NP-complete, so we can prove this probelem X is NP-complete by proving $Y \leq_p X$.

For any given positive integers $x_1, ..., x_n$, and B, we can let $k = 2$ and $\Sigma x_i = S$. As subset sum B problem is equivalent to subset sum $S - B$ problem, we assume $B \geq \frac{S}{2}$.

If $B = \frac{S}{2}$ and there exists a partition satisfying: $\sum_{i=1}^{k} \left(\sum_{x_j \in S_i} x_j \right)^2 \leq \frac{S^2}{2}$, otherwise, if $B > \frac{S}{2}$ there exists a partition satisfying:

$$(B-1)^2 + (S - B + 1)^2 < \sum_{i=1}^{k} \left(\sum_{x_j \in S_i} x_j \right)^2 \leq B^2 + (S - B)^2$$

The subset sum of this partition is exactly B.

We know by cauchy's inequality $\frac{S^2}{2} \leq x^2 + (S - x)^2 \leq S^2$, so, it takes no more than $O(log(S^2))$ calls of problem Y to know wether problem X can be sovled. So $Y \leq_p X$ and this problem is NP-complete.

p653 Ex11.5

Answer

The optimal makespan is at least:

$$T^* \geq \frac{1}{m} \sum_k t_k$$

For client which finished its assigned jobs at the latest (T_j), when we assigned the last job t_i to it:

$$T_j - t_i \leq \frac{1}{m} \sum_k t_k \leq T^*$$

$$T_j = (T_j - t_i) + t_i$$

As the sizes of jobs range between 1 and 50, and the total load is always at least 3,000, $t_i \leq 50 \div (3000 \div 10)T^* \leq 20\% T^*$

So, the Greedy-Balance Algorithm will always find a solution whose makespan is at most 20 percent above the average load.

p654 Ex11.6

Answer

The optimal makespan is at least:

$$T^* \geq \frac{\sum_l t_l}{2k + m}$$

Algorithm 15 Algorithm of Ex11.6

1: Stat with no jobs assigned

2: Set $T_i = 0$ and $A(i) = \emptyset$ for all macines M_i

3: **for** j=1,...n **do**

4:     Let M_i be a machine that achieves the minimum $min_k T_k$

5:     Assign job j to machine M_i

6:     Set $A(i) = A(i) \cup \{j\}$

7:     Set $T_j = T_j + (t_j$ or $\frac{1}{2}t_j)$(depend on M_j slow or fast)

8: **end for**

For client which finished its assigned jobs at the latest (T_j), when we assigned the last job t_i to it:

$$T_j - t_i \leq \frac{\sum_l t_l}{2k + m} \leq T^*$$

We also have:

$$T^* \geq \frac{1}{2}t_i$$

so,

$$T_j = (T_j - t_i) + t_i \leq T^* + 2T^* = 3T^*$$

The algorithm is polynomial-time.

Homework 11

p653 Ex11.4

Consider an optimization version of the Hitting Set Problem defined as follows. We are given a set $A = \{a_1, \ldots, a_n\}$ and a collection $B_1, B_2, \ldots, B_m$ of subsets of A. Also, each element $a_i \in A$ has a weight $w_i \geq 0$. The problem is to find a hitting set $H \subseteq A$ such that the total weight of the elements in H, that is, $\sum_{a_i \in H} w_i$, is as small as possible. (As in Exercise 5 in Chapter 8, we say that H is a hitting set if $H \cap B_i$ is not empty for each i..) Let $b = \max_i |B_i|$ denote the maximum size of any of the sets $B_1, B_2, \ldots, B_m$. Give a polynomial-time approximation algorithm for this problem that finds a hitting set whose total weight is at most b times the minimum possible.

Answer

Consider the following problem for Linear Programming Relaxation.

$$\text{Min} \sum_{i=1}^{n} w_i x_i$$

$$\text{s.t.} \quad 0 \leq x_i \leq 1 \qquad \text{for all } i = 1, \ldots, n$$

$$\sum_{i:a_i \in B_j} x_i \geq 1 \qquad \text{for all } j = 1, \ldots, m$$

After we get the relaxation solution $R = x_i$ of this problem, we can round x to get the rounding solution:

$$S = \{a_i \mid x_i \geq 1/b\}$$

First of all, S is a hitting set, because for any B_j, there is at least one $a_k \in B_j, a_k \geq 1/b$. Secondly, the relaxation solution is no worse than the original problem. Last, the rounding solution is at least b times of relaxation solution. So:

$$w(S) \leq bw(R) \leq bw(S^*)$$

p655 Ex11.7

Answer

(a)

This problem is similar to load balancing problem. We can use Greedy-Balance algorithm: The optimal spread is at most:

$$S^* \leq \frac{\sum_j v_j}{m}$$

Algorithm 16 Algorithm of Ex11.7

1: Stat with no vister is showed ad.

2: Set $V_i = 0$ and $C_i = \emptyset$ for all ads A_i

3: **for** j=1,...n **do**

4:     Let M_i be the ad that achieves the minimum $min_k V_k$

5:     Show ad M_i to visitor j.

6:     Set $C_i = C_i \cup \{j\}$

7:     Set $V_i = V_j + v_j$

8: **end for**

9: Let S be the ad that achieves the minimum $min_k V_k$

10: **return** S

---

and the ad achieve most V_K is larger than :

$$V_{max} \geq \frac{\sum_j v_j}{m}$$

If $S < \frac{\sum_j v_j}{2m}$. For ad k which achieved (V_{max}), when the last time we show it to visitor i, as $v_i \leq \frac{\sum_j v_j}{2m}$:

$$V_{max} - v_i \geq \frac{\sum_j v_j}{2m} \geq S^*$$

This is contraydictory to the algorithm above, so:

$$S^* \leq \frac{\sum_j v_j}{m} \leq 2S$$

(a)

Suppose there are 13 customers with values 1,1,1, 1,1,1, 1,1,1, 1,1,1, 4. The optimal solution is 4, in this sample instance, but my algorithm finds 3.

p655 Ex11.8

Answer

The conjecture is true. We can retract from the optimal solution. we choose load i, which is the last task of machine j and without which machine j finish loads first as the last comming load. Repeat this to find the second from the bottom until the first load. So the sequence is found.

Although this is always possible, it is not known wether it can be solved in polynomial time.

p656 Ex11.9

Consider the following maximization version of the 3 -Dimensional Matching Problem. Given disjoint sets X, Y, and Z, and given a set $T \subseteq X \times Y \times Z$ of ordered triples, a subset $M \subseteq T$ is a 3 -dimensional matching if each element of $X \cup Y \cup Z$ is contained in at most one of these triples. The Maximum 3-Dimensional Matching Problem is to find a 3 -dimensional matching M of maximum size. (The size of the matching, as usual, is the number of triples it contains. You may assume $|X| = |Y| = |Z|$ if you want.) Give a polynomial-time algorithm that finds a 3 -dimensional matching of size at least $\frac{1}{3}$ times the maximum possible size.

Answer

We can choose $t_i \in T$ one by one until we cannot choose one more t_i. First of all, any triple in optimal solution has at least one element belows to this approximation solution. Otherwise, we can add this triple to my solution. For any $t_i = (x_i, y_i, z_i)$ in my solution, x_i, y_i, z_i can belows to at most three triple. So my algorithm finds a 3 -dimensional matching of size at least $\frac{1}{3}$ times the maximum possible size.

Homework 12

p656 Ex11.10

Answer

(a)

By proof by contradiction, if there exists a node $v \in T$ but $v \notin S$. according to the algorithm, there must be some of its neighbors $v' \in S$. If all these neighbors $w(v') < w(v)$, before pick first node of these neighbors, v is bigger than all of them, according to the algorithm, v should be picked. That is contraydictory to the assumption that $v \notin S$.

(b)

Let $N(v)$ be all neighbors of v. For any $v' \in S, v' \notin T$, v' could be neighbor of at most 4 nodes $v \in T$, so:

$$\sum_{v \in T} w(v) = \sum_{v \in T, v \in S} w(v) \quad + \sum_{v \in T, v \notin S} w(v)$$

$$\leq \sum_{v \in T, v \in S} w(v) \quad + \sum_{v \in T, v \notin S} \sum_{v' \in N(v)} w(v')$$

$$\leq \sum_{v \in S, v \in T} w(v) \quad + 4 \sum_{v \in S, v \notin T} w(v)$$

$$\leq 4 \sum_{v \in S} w(v)$$

p656 Ex11.10

Answer

We can round all weights to $w_i' = [\frac{n}{2\epsilon W} w_i]$ and fix the problem with the limitation of $W' = \frac{n}{2\epsilon}$. Let S be the solultion:

$$\sum_{i \in S} w_i \leq \sum_{i \in S} (w_i' + 0.5) \frac{2\epsilon W}{n} \leq (\frac{n}{2\epsilon} + 0.5n) \frac{2\epsilon W}{n} \leq (1 + \epsilon)W$$

p703 Ex12.2

Answer

(a)

$$G = \{v_1, v_2, u_1, u_2, v_1 - u_1, v_1 - u_2, v_2 - u_2\}$$

If we add $v_1 - u_2$ to matching at first, then this gradient ascent algorithm does not return the maximum matching$\{v_1 - u_1, v_2 - u_2\}$

(b)

By proof by contradiction, if for any $e' \in M'$, $M \cup e'$ is not a matching in G, then for these $e_i' = v_i - u_i$, $v_i \in M$ or $u_i \in M$. So there is at least $|M'|$ different nodes in M, that is, $|M| \geq \frac{1}{2}|M'|$, which is contraydictory to the fact that $2|M| < |M'|$

(c)

By the above Gradient Ascent Algorithm, we can always add edge to M until for any matching M', $|M| \geq \frac{1}{2}|M'|$. So, any locally optimal matching returned by the gradient ascent algorithm in a bipartite graph G is at least half as large as a maximum matching in G.

p703 Ex12.3

Answer

(a)

If $T_1 > 2T_2$ or $2T_1 > T_2$, let's say that $T_1 \geq 2T_2$, there must exists $t_i \in T_1$, $t_i \leq \frac{1}{2}(T_1 + T_2)$. Then we can move t_i to T_2 and make the gap smaller.

(b)

Each time when we move t_j to other machine, t_j is no more than the gap. So, at the next move, available t_j' is samller than the gap so is smaller than t_j. That means under this variant, each job will move at most once.

(c)

Machine1:{3,4}, Machine2:{4,5}.

Homework 13

p705 Ex12.4

Answer

(a)

We can define
$$\Delta = \sum_{i \neq j} |T_i - T_j|$$
After each swap, For $T_k < min\{T_i, T_j\}$ and $T_k > min\{T_i, T_j\}$ their total distance from T_i and T_j is the same. For $min\{T_i, T_j\} < T_k < max\{T_i, T_j\}$, their total distance from T_i and T_j is smaller. As $max\{T_i', T_j'\} < max\}T_i, T_j\}$, so:
$$0 \leq \Delta' = \sum_{i \neq j} |T_i' - T_j'| < \Delta$$

so, the local search heuristic terminates in a finite number of steps, with a stable assignment, on any instance.

(b)

It is obvious that the maximum t_i is not more than minimum possible makespan. If the stable assignment has a makespan that is larger than factor of 2 of minimum possible makespan. Assume i is the machine which have the most load in stable assignment. There exist a $B(i) \subseteq A(i)$, $T_{optimal} < T_{B(i)} < T_{A(i)}$ Then we can swap $B(i)$ with any $A(j) < T_{optimal}$. This is contraydictory to stable assignment.

p782 Ex13.1

Answer

we can color each node randomly. For each edge, it has a probility of $\frac{A_3^2}{3*3} = \frac{2}{3}$ to be satisfied. So,
$$exp(|E_{satisfied}|) = \frac{2}{3}|E| \geq \frac{2}{3}c^*$$

p783 Ex13.3

Answer

(a)

    It is abvious that this protocol is conflict-free because P_i enter the set S only if each of the processes with which it is in conflict chooses the value 0, that means all of these processes didn't enter S.

$$E(|S|) = n * (\frac{1}{2})^{d+1}$$

(b)

$$E(|S|) = n * p(1-p)^d$$

$$\frac{dy}{dp}E(|S|) = n[(1-p)^d - pd(1-p)^{d-1}] = 0$$

So, $p = \frac{1}{d+1}$, and

$$E(|S|) = \frac{nd^d}{(d+1)^{d+1}}$$